

# Deep Classification and Semantic Segmentation of 3D Point cloud

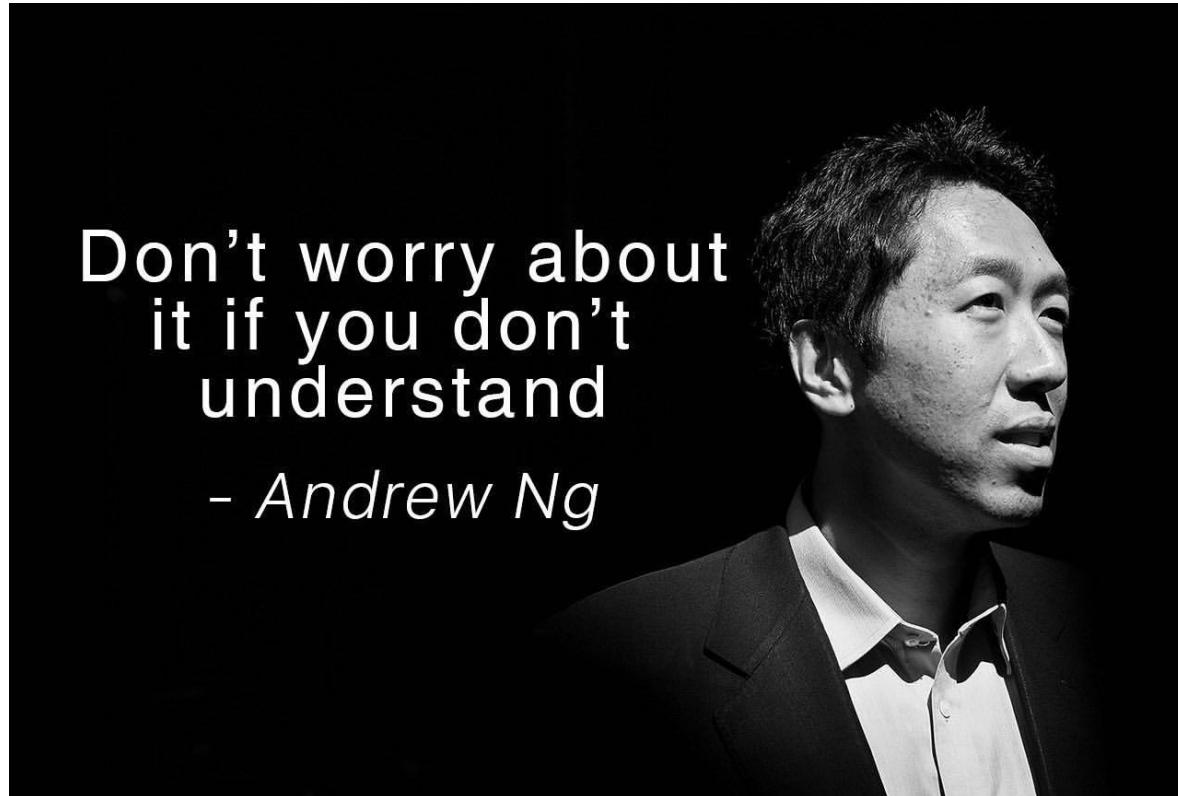
Maryam Jameela

Workshop: Deep Learning

**Date:** June 21, 2021 13:00 pm to 15:00 pm Mountain Daylight Time

# Agenda

- System Configuration
- Classification
  - PointNet
  - ModelNet40
- Segmentation
  - 3D U-Net
  - Utility Data

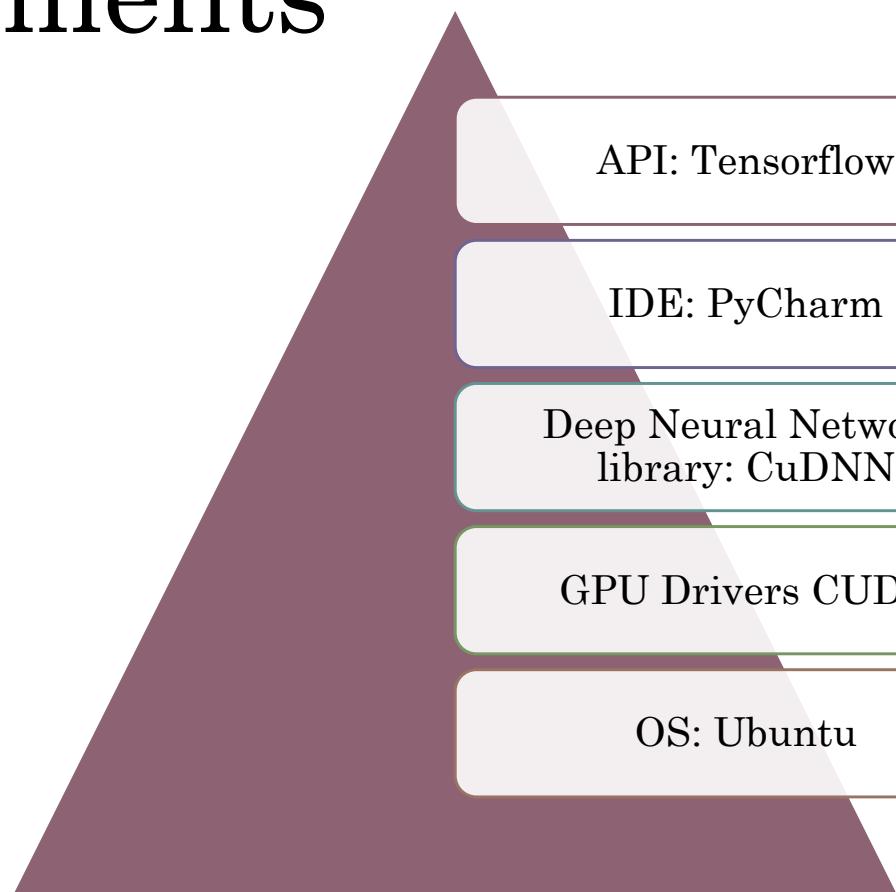


Don't worry about  
it if you don't  
understand  
*- Andrew Ng*

# System Configuration

# System Requirements

- Operating System
  - Ubuntu 18.04
- Environment
  - CUDA 10.0
  - CuDNN 7.4
  - Pycharm 2021
- Install
  - Tensorflow 1.14
  - Numpy
  - H5py



API: Tensorflow

IDE: PyCharm

Deep Neural Network  
library: CuDNN

GPU Drivers CUDA

OS: Ubuntu

# Compatibility Table

Version	Python version	Compiler	Build tools	cuDNN	CUDA
tensorflow-2.5.0	3.6-3.9	GCC 7.3.1	Bazel 3.7.2	8.1	11.2
tensorflow-2.4.0	3.6-3.8	GCC 7.3.1	Bazel 3.1.0	8.0	11.0
tensorflow-2.3.0	3.5-3.8	GCC 7.3.1	Bazel 3.1.0	7.6	10.1
tensorflow-2.2.0	3.5-3.8	GCC 7.3.1	Bazel 2.0.0	7.6	10.1
tensorflow-2.1.0	2.7, 3.5-3.7	GCC 7.3.1	Bazel 0.27.1	7.6	10.1
tensorflow-2.0.0	2.7, 3.3-3.7	GCC 7.3.1	Bazel 0.26.1	7.4	10.0
tensorflow_gpu-1.15.0	2.7, 3.3-3.7	GCC 7.3.1	Bazel 0.26.1	7.4	10.0
tensorflow_gpu-1.14.0	2.7, 3.3-3.7	GCC 4.8	Bazel 0.24.1	7.4	10.0
tensorflow_gpu-1.13.1	2.7, 3.3-3.7	GCC 4.8	Bazel 0.19.2	7.4	10.0

# CUDA

## CUDA Toolkit 10.0 Archive

### Select Target Platform

Click on the green buttons that describe your target platform. Only supported platforms will be shown.

#### Operating System

Windows    Linux    Mac OSX

#### Architecture

x86\_64    ppc64le

#### Distribution

Fedora    OpenSUSE    RHEL    CentOS    SLES    Ubuntu

#### Version

18.04    16.04    14.04

#### Installer Type

runfile {local}    deb {local}    deb {network}    cluster {local}

### Download Installers for Linux Ubuntu 18.04 x86\_64

The base installer is available for download below.

There is 1 patch available. This patch requires the base installer to be installed first.

#### Base Installer

Download (1.6 GB) 

#### Installation Instructions:

1. `sudo dpkg -i cuda-repo-ubuntu1804-10-0-local-10.0.130-410.48\_1.0-1\_amd64.deb`
2. `sudo apt-key add /var/cuda-repo-<version>/7fa2af80.pub`
3. `sudo apt-get update`
4. `sudo apt-get install cuda`

Other installation options are available in the form of meta-packages. For example, to install all the library packages, replace "cuda" with the "cuda-libraries-10-0" meta package.

For more information on all the available meta packages click [here](#).

# CuDNN

1. Register a nvidia developer account and [download cudnn](#)
2. Check where your cuda installation is.
3. Copy the files:

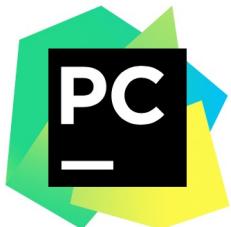
```
$ cd folder/extracted/contents  
$ sudo cp include/cudnn.h /usr/local/cuda/include  
$ sudo cp lib64/libcudnn* /usr/local/cuda/lib64  
$ sudo chmod a+r /usr/local/cuda/lib64/libcudnn*
```

# PyCharm

PyCharm

Coming in 2021.2 What's New Features Learn Buy

Download



Version: 2021.1.2

Build: 211.7442.45

1 June 2021

[System requirements](#)

[Installation Instructions](#)

[Other versions](#)

## Download PyCharm

Windows

macOS

Linux

### Professional

For both Scientific and Web Python development. With HTML, JS, and SQL support.

[Download](#)

[Free trial](#)

### Community

For pure Python development

[Download](#)

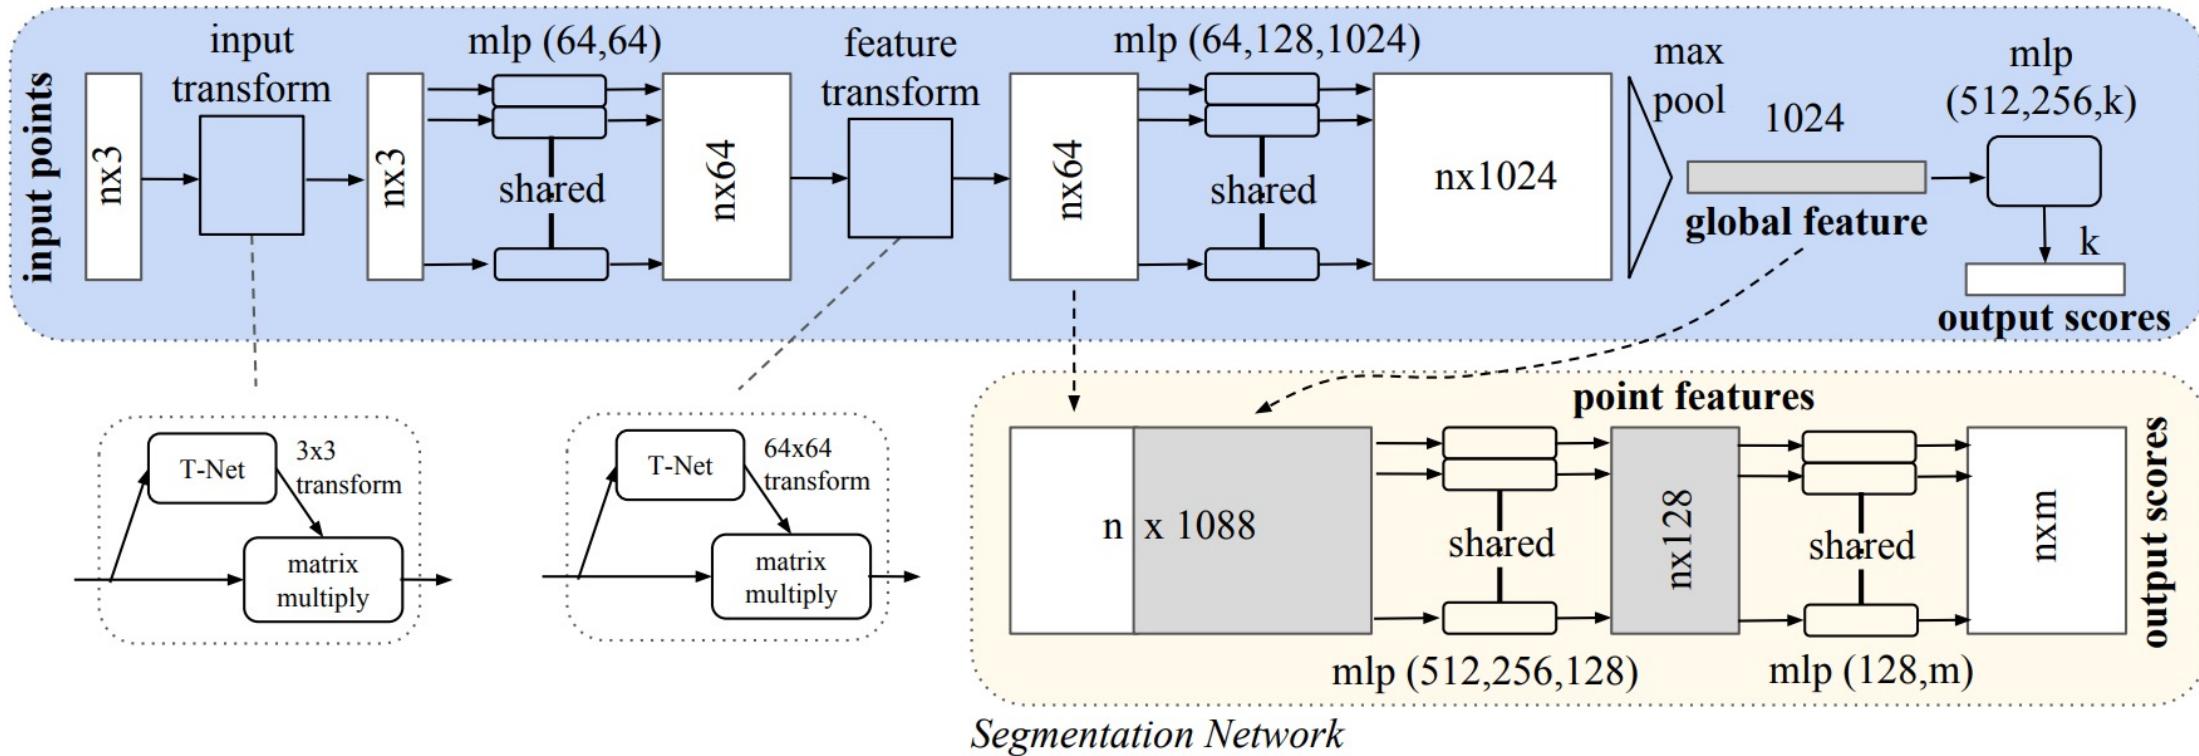
[Free, open-source](#)

# Tensorflow

- pip3 install tensorflow==1.14.0
- pip3 install tensorflow-gpu==1.14.0
- pip3 install numpy
- pip3 install h5py

# Demo

## Classification Network

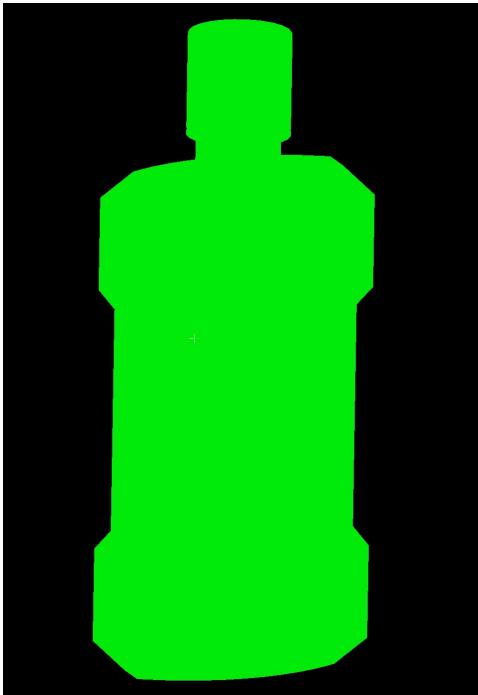


# PointNet: Classification Demo

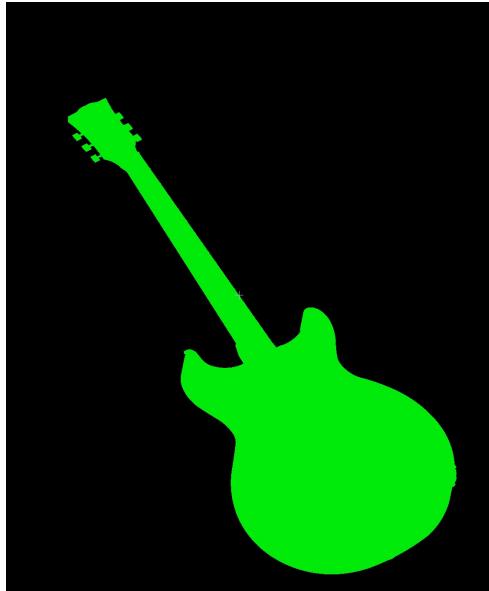
Pointnet: Deep Learning on Point Sets for 3d Classification and Segmentation [Qi et al., 2017]

# Dataset

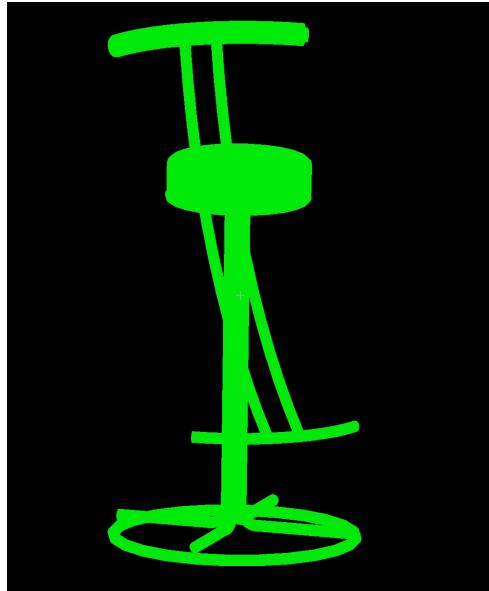
- ModelNet40: 40 Classes
- Examples:



Bottle



Guitar

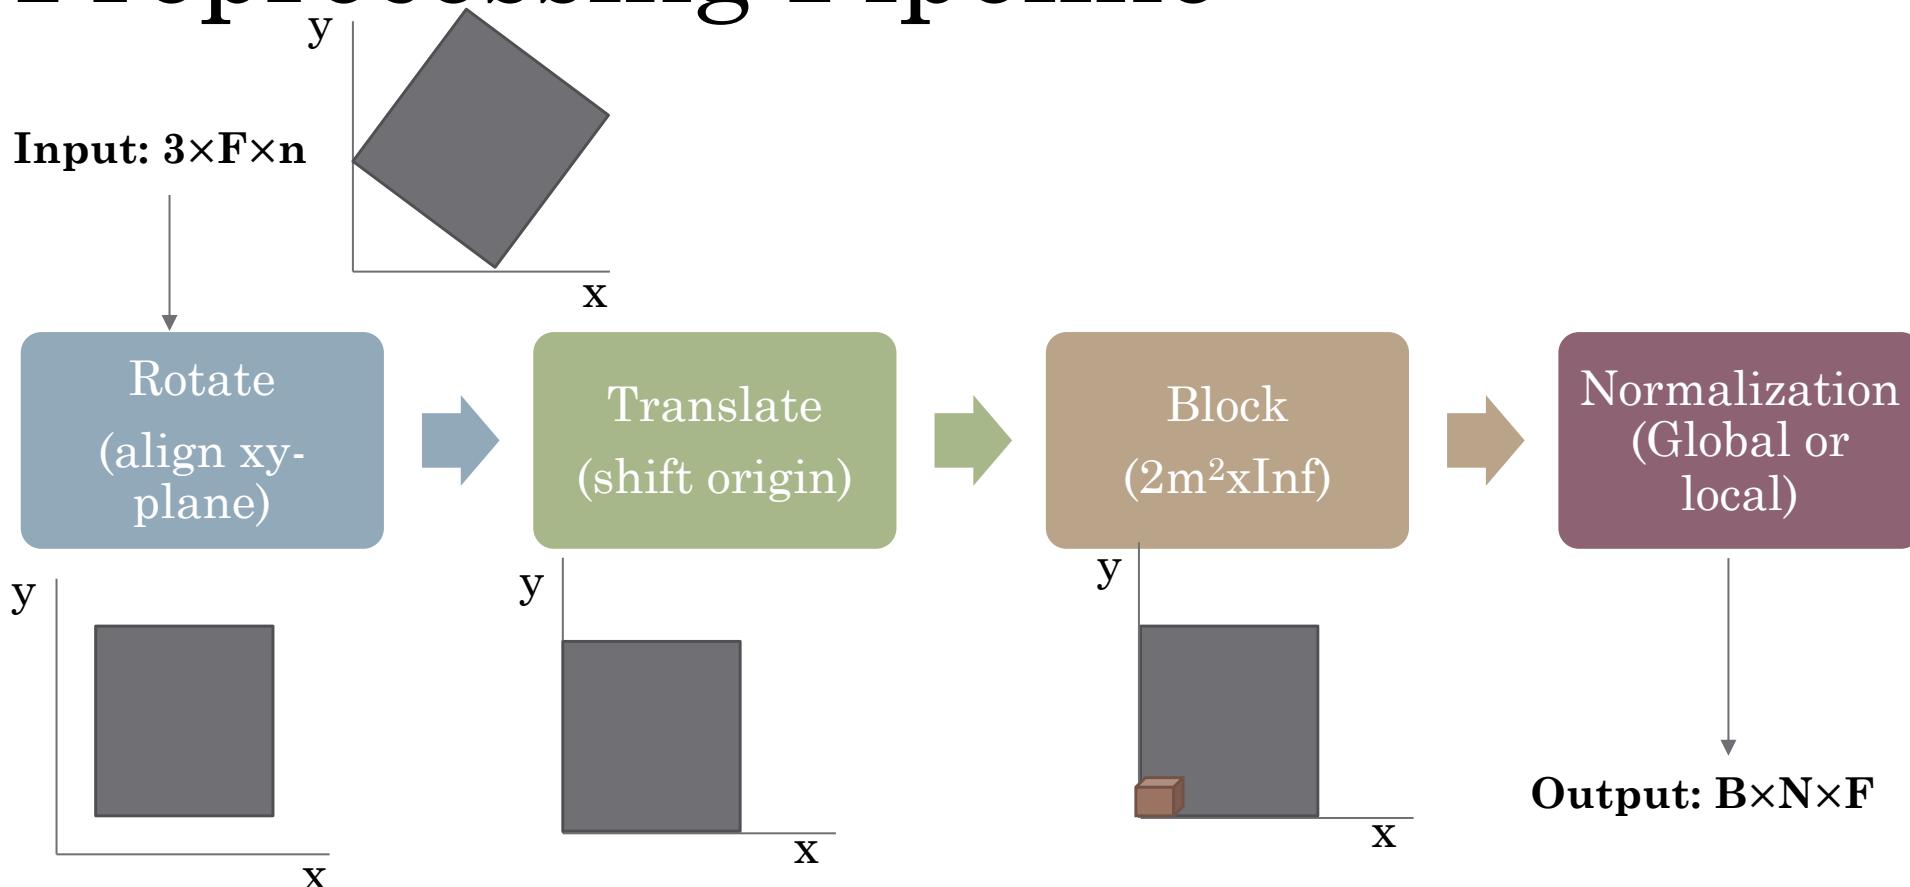


Stool



Airplane

# Preprocessing Pipeline



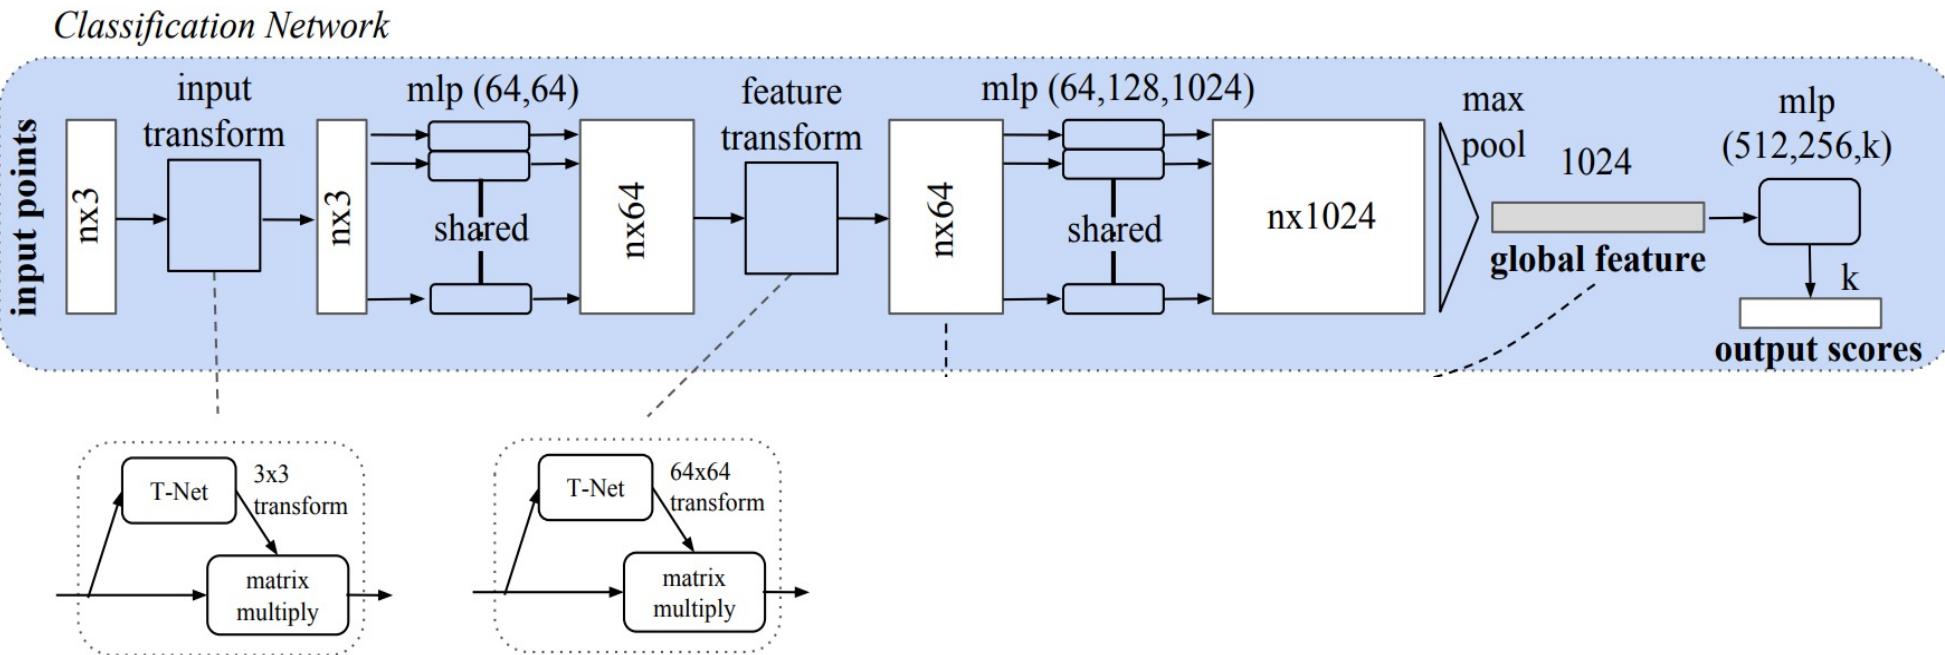
F: Features

B: Blocks

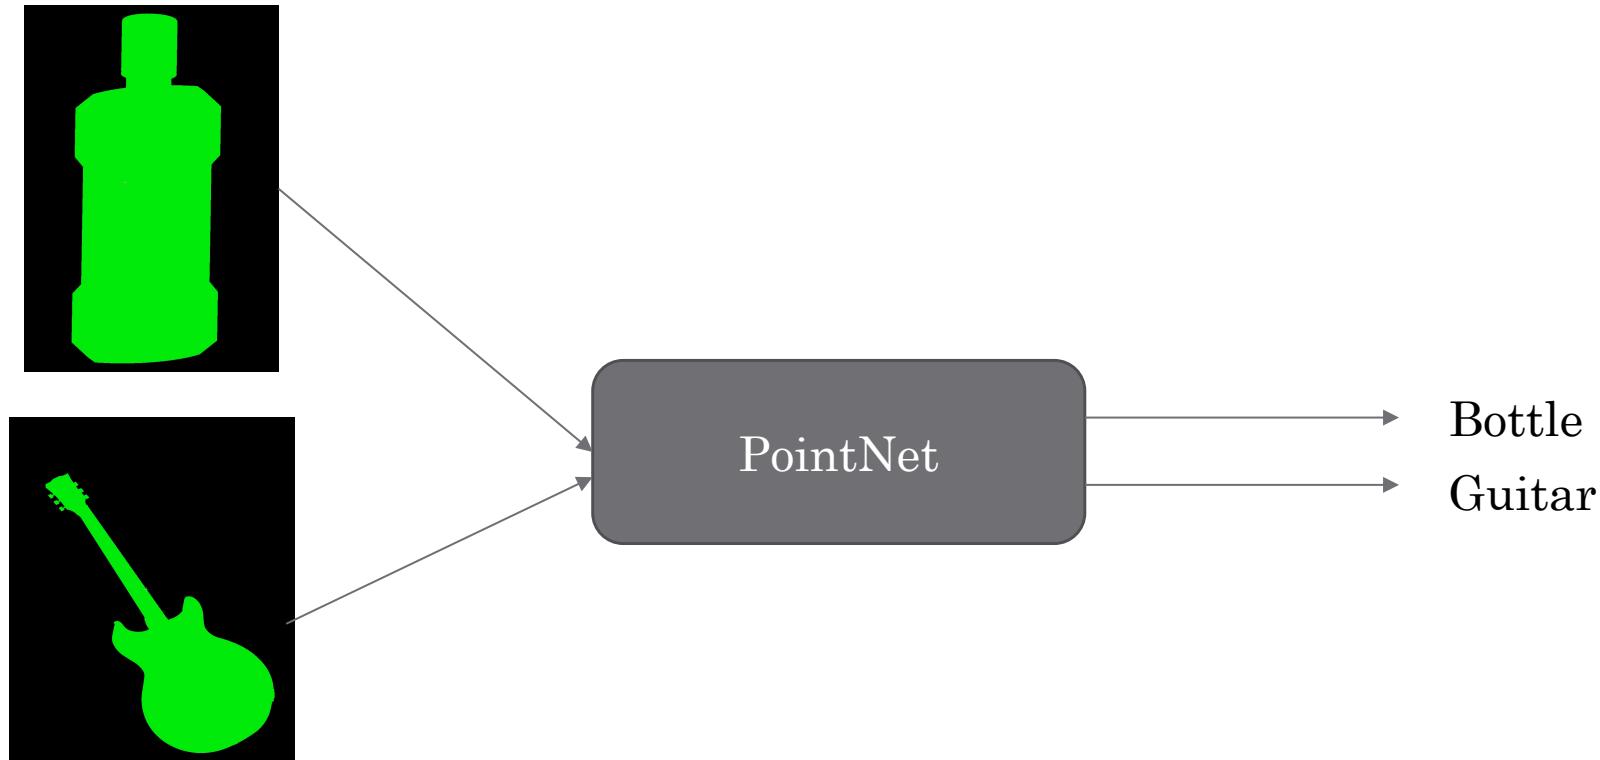
N: Maximum number of points in a block

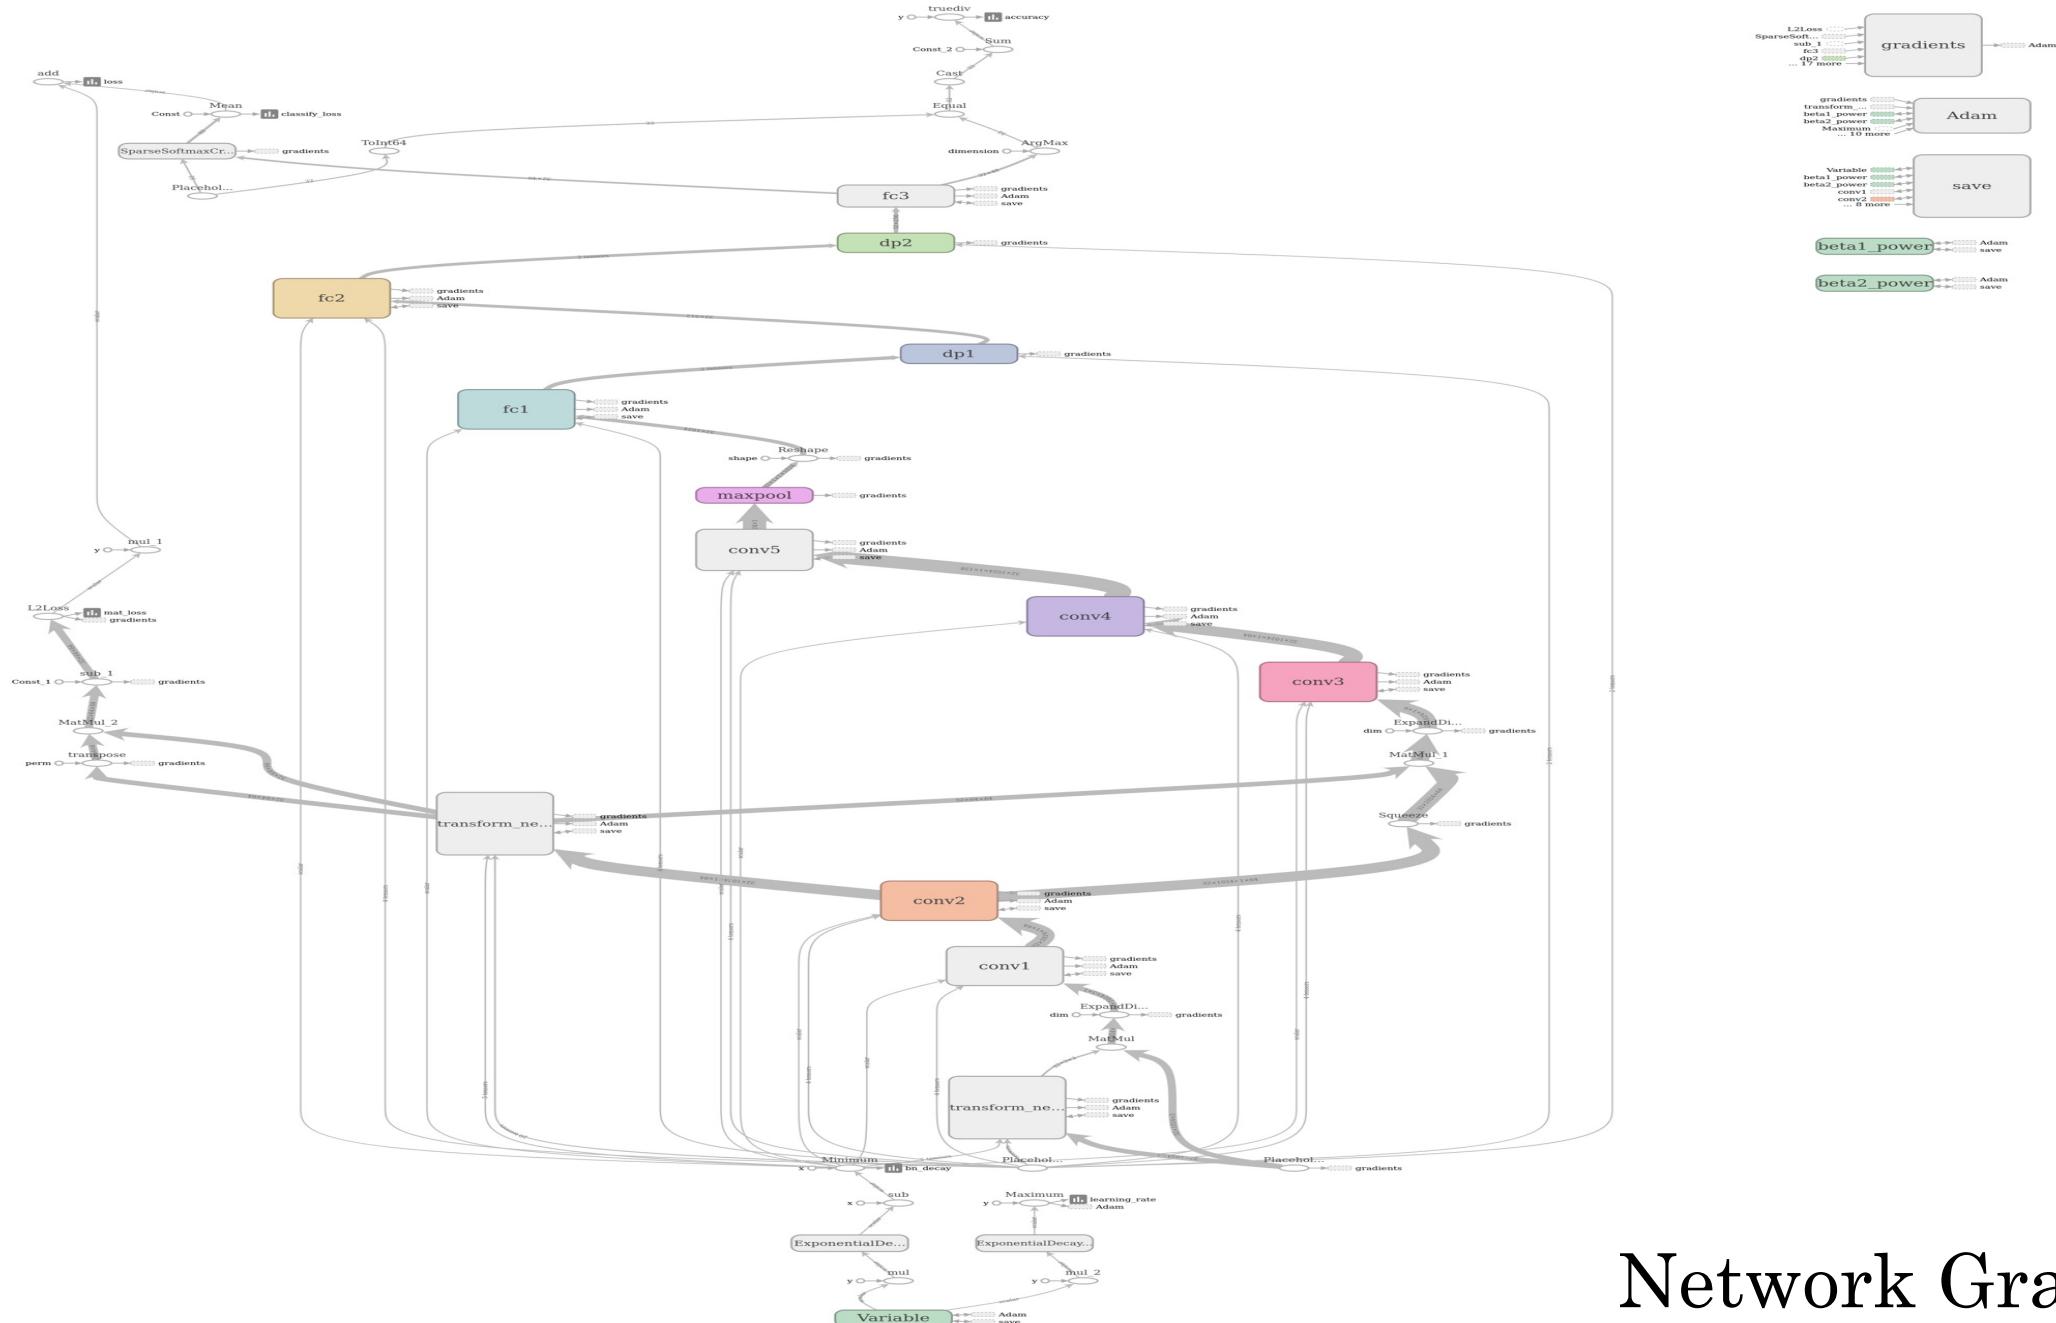
n: Total number of points in point cloud

# Network



# Classification

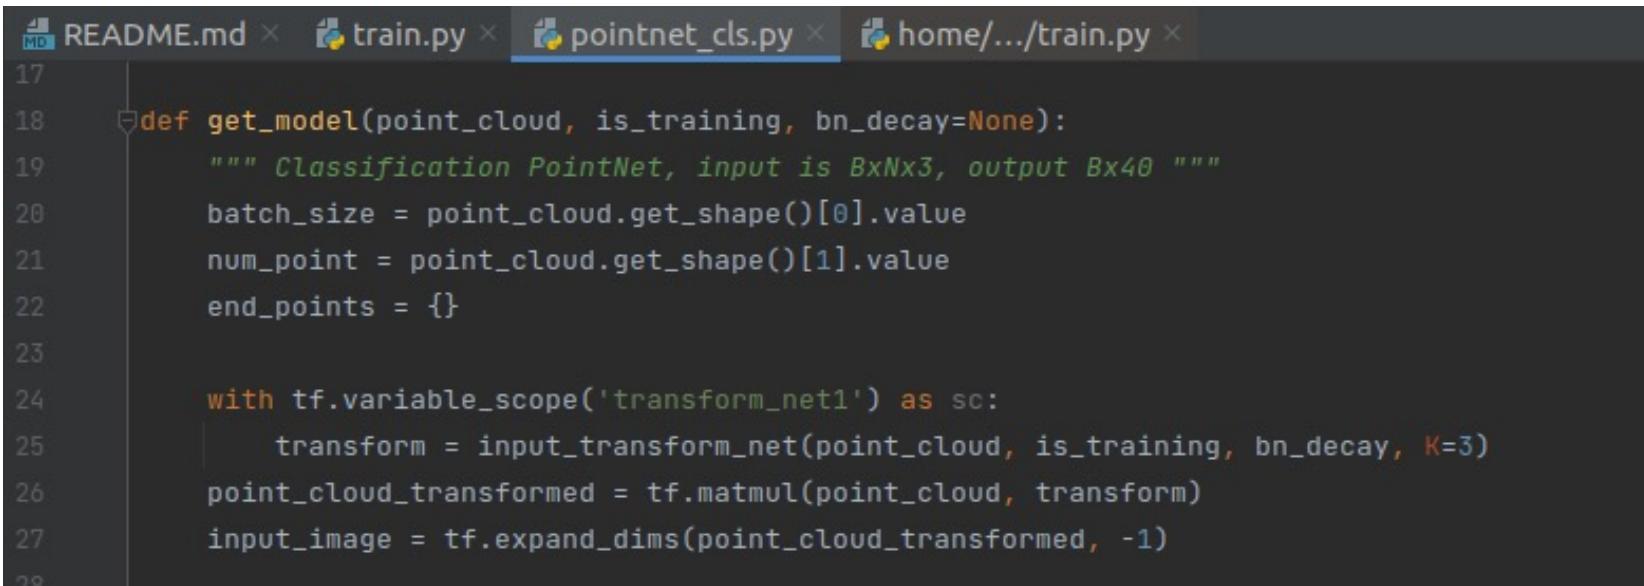




# Network Graph

# Input T-Net

First Spatial Transformer is to adjust the point cloud in the space. Intuitively, rotating an angle, such as turning an object to the front, is conducive to classification or segmentation.



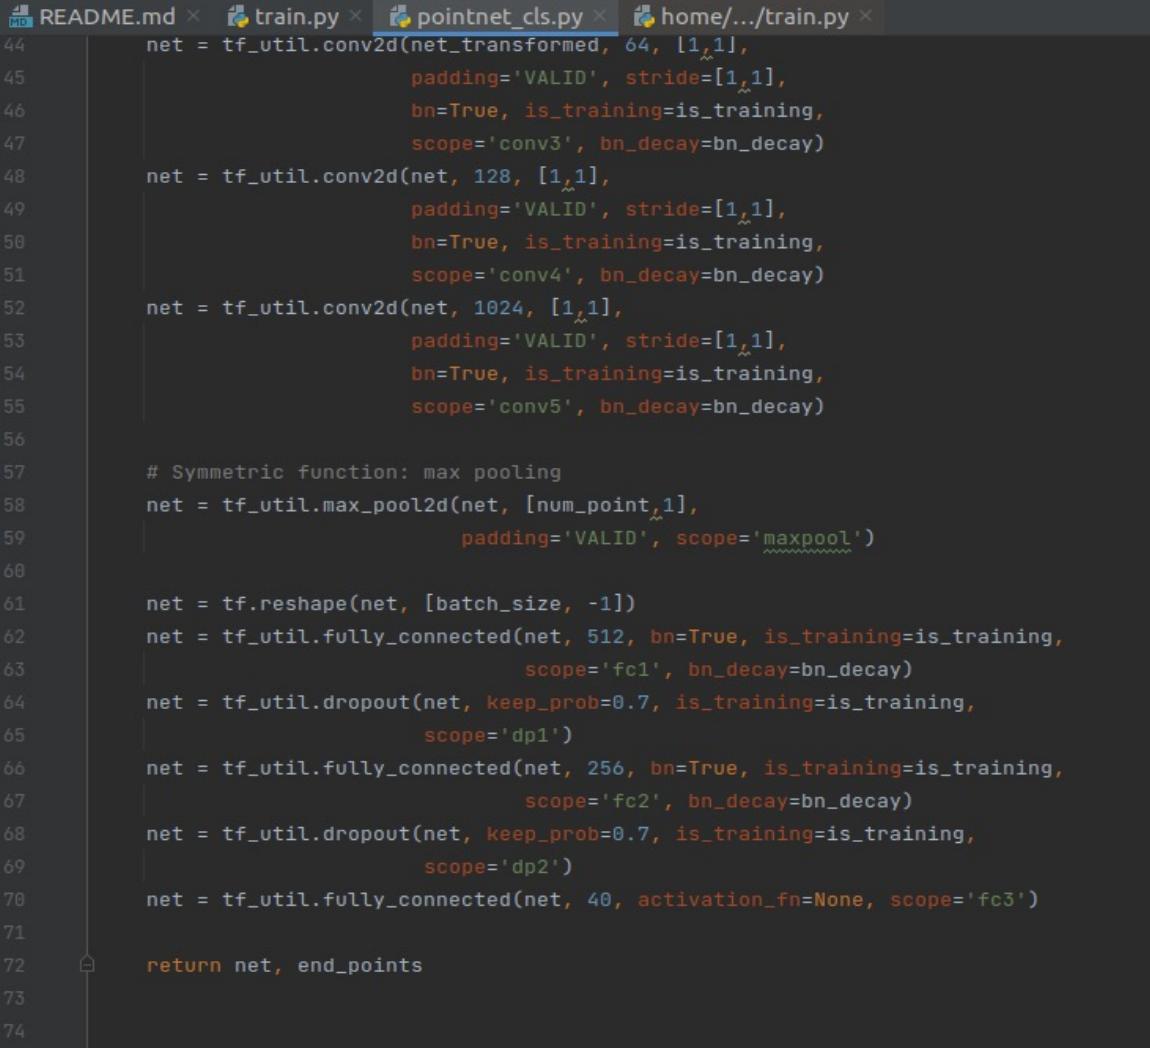
```
17
18     def get_model(point_cloud, is_training, bn_decay=None):
19         """ Classification PointNet, input is BxNx3, output Bx40 """
20         batch_size = point_cloud.get_shape()[0].value
21         num_point = point_cloud.get_shape()[1].value
22         end_points = {}
23
24         with tf.variable_scope('transform_net1') as sc:
25             transform = input_transform_net(point_cloud, is_training, bn_decay, K=3)
26             point_cloud_transformed = tf.matmul(point_cloud, transform)
27             input_image = tf.expand_dims(point_cloud_transformed, -1)
28
```

# Feature T-Net

The feature transformer is to convert the features into high semantic information of the scene

```
18     def get_model(point_cloud, is_training, bn_decay=None):
19         """ Classification PointNet, input is BxNx3, output Bx40 """
20         batch_size = point_cloud.get_shape()[0].value
21         num_point = point_cloud.get_shape()[1].value
22         end_points = {}
23
24         with tf.variable_scope('transform_net1') as sc:
25             transform = input_transform_net(point_cloud, is_training, bn_decay, K=3)
26             point_cloud_transformed = tf.matmul(point_cloud, transform)
27             input_image = tf.expand_dims(point_cloud_transformed, -1)
28
29             net = tf_util.conv2d(input_image, 64, [1,3],
30                                 padding='VALID', stride=[1,1],
31                                 bn=True, is_training=is_training,
32                                 scope='conv1', bn_decay=bn_decay)
33             net = tf_util.conv2d(net, 64, [1,1],
34                                 padding='VALID', stride=[1,1],
35                                 bn=True, is_training=is_training,
36                                 scope='conv2', bn_decay=bn_decay)
37
38         with tf.variable_scope('transform_net2') as sc:
39             transform = feature_transform_net(net, is_training, bn_decay, K=64)
40             end_points['transform'] = transform
41             net_transformed = tf.matmul(tf.squeeze(net, axis=[2]), transform)
42             net_transformed = tf.expand_dims(net_transformed, [2])
43
```

# Global Feature, Dropout & Classification



A screenshot of a code editor showing a Python script named `pointnet_cls.py`. The script contains code for a neural network architecture, specifically for global feature extraction and classification. The code includes several layers of convolutional and fully connected layers, along with dropout and max pooling operations. The script is part of a larger project, as indicated by the tabs at the top of the editor.

```
44     net = tf_util.conv2d(net_transformed, 64, [1,1],
45                         padding='VALID', stride=[1,1],
46                         bn=True, is_training=is_training,
47                         scope='conv3', bn_decay=bn_decay)
48     net = tf_util.conv2d(net, 128, [1,1],
49                         padding='VALID', stride=[1,1],
50                         bn=True, is_training=is_training,
51                         scope='conv4', bn_decay=bn_decay)
52     net = tf_util.conv2d(net, 1024, [1,1],
53                         padding='VALID', stride=[1,1],
54                         bn=True, is_training=is_training,
55                         scope='conv5', bn_decay=bn_decay)
56
57     # Symmetric function: max pooling
58     net = tf_util.max_pool2d(net, [num_point,1],
59                             padding='VALID', scope='maxpool')
60
61     net = tf.reshape(net, [batch_size, -1])
62     net = tf_util.fully_connected(net, 512, bn=True, is_training=is_training,
63                                  scope='fc1', bn_decay=bn_decay)
64     net = tf_util.dropout(net, keep_prob=0.7, is_training=is_training,
65                          scope='dp1')
66     net = tf_util.fully_connected(net, 256, bn=True, is_training=is_training,
67                                  scope='fc2', bn_decay=bn_decay)
68     net = tf_util.dropout(net, keep_prob=0.7, is_training=is_training,
69                          scope='dp2')
70     net = tf_util.fully_connected(net, 40, activation_fn=None, scope='fc3')
71
72     return net, end_points
73
74
```

# Experimental Configuration

- Configuration
  - 70-30 Split Training and Test
  - Block size=  $2m^2 \times Inf$ ,
  - Batch size=32x2048x6
  - 1 GPU GTX 2080
  - Training Time: 12 hours (250 Epochs)
  - Inference: 30 seconds
- Evaluation Configuration
  - Accuracy

# Training

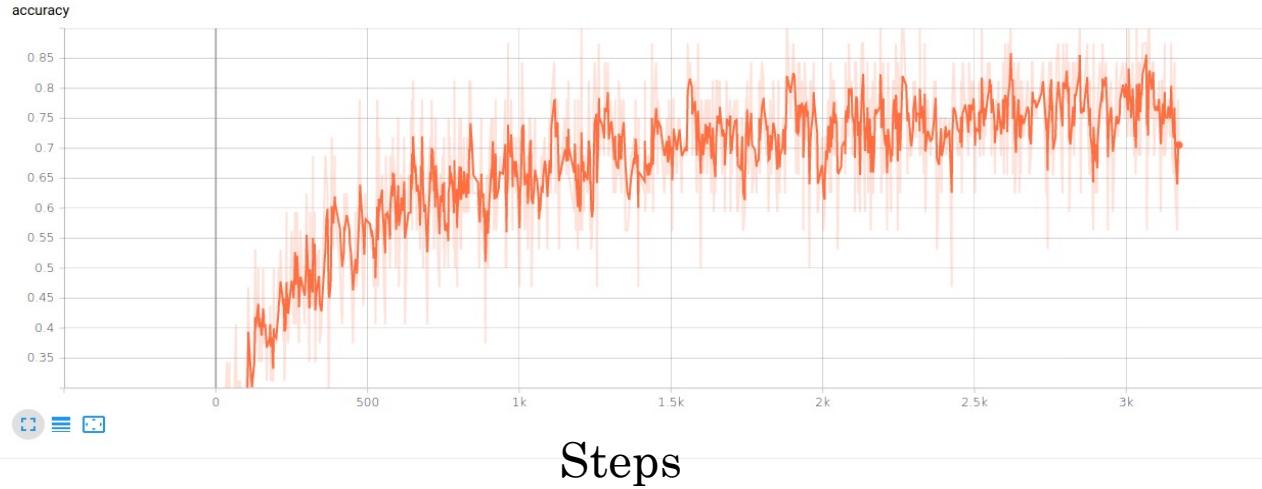
- python3 train.py
- Hyperparameter:

```
parser = argparse.ArgumentParser()
parser.add_argument('--gpu', type=int, default=0, help='GPU to use [default: GPU 0]')
parser.add_argument('--model', default='pointnet_cls', help='Model name: pointnet_cls or pointnet_cls_basic [default: pointnet_cls]')
parser.add_argument('--log_dir', default='log', help='Log dir [default: log]')
parser.add_argument('--num_point', type=int, default=1024, help='Point Number [256/512/1024/2048] [default: 1024]')
parser.add_argument('--max_epoch', type=int, default=250, help='Epoch to run [default: 250]')
parser.add_argument('--batch_size', type=int, default=32, help='Batch Size during training [default: 32]')
parser.add_argument('--learning_rate', type=float, default=0.001, help='Initial learning rate [default: 0.001]')
parser.add_argument('--momentum', type=float, default=0.9, help='Initial learning rate [default: 0.9]')
parser.add_argument('--optimizer', default='adam', help='adam or momentum [default: adam]')
parser.add_argument('--decay_step', type=int, default=200000, help='Decay step for lr decay [default: 200000]')
parser.add_argument('--decay_rate', type=float, default=0.7, help='Decay rate for lr decay [default: 0.8]')
FLAGS = parser.parse_args()
```

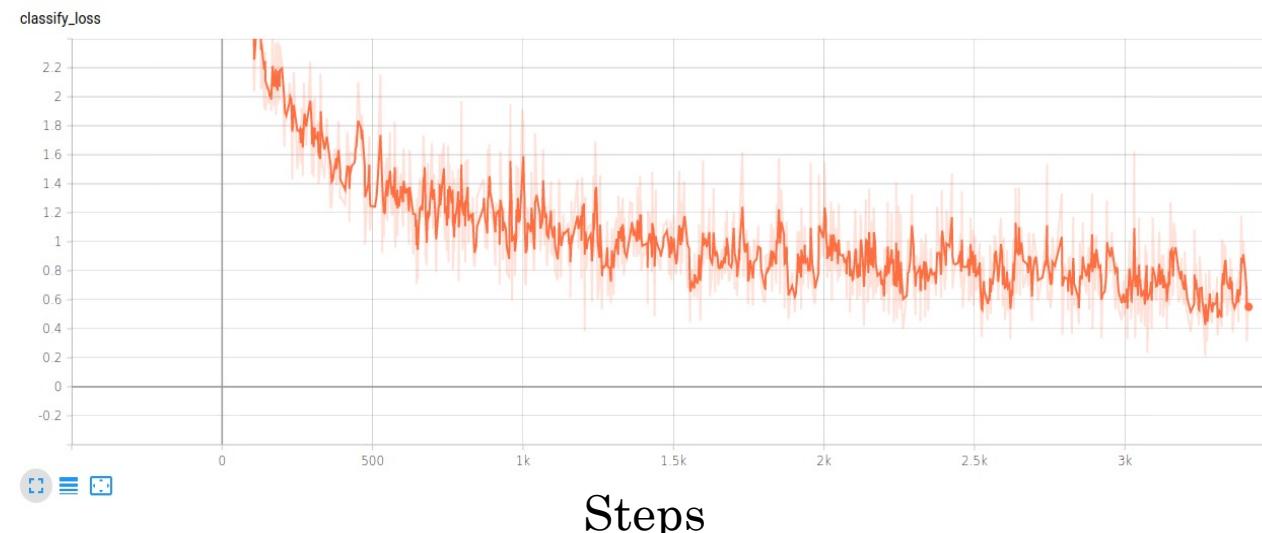
- tensorboard --logdir log

# Convergence

Accuracy

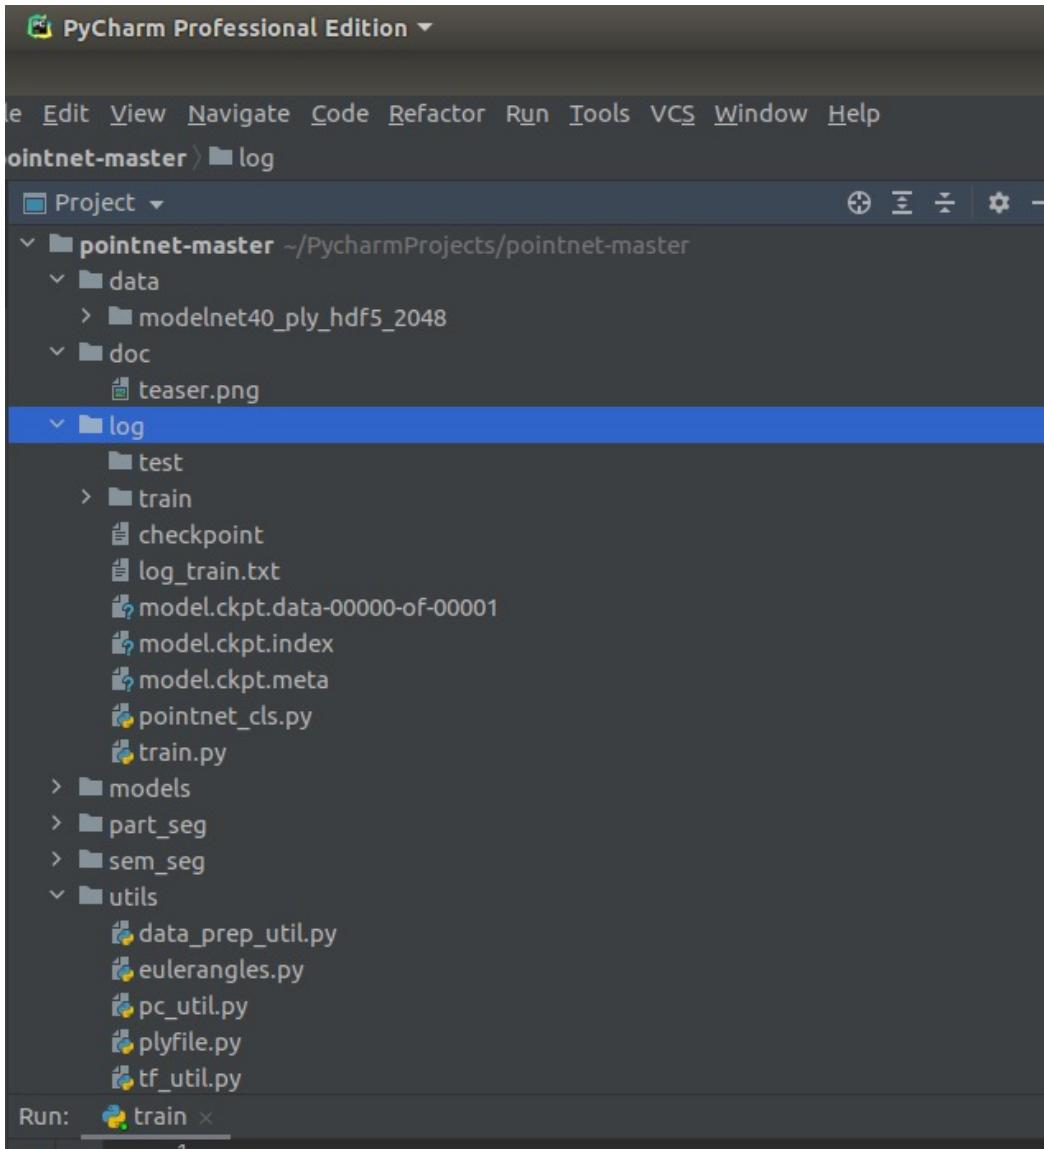


Loss



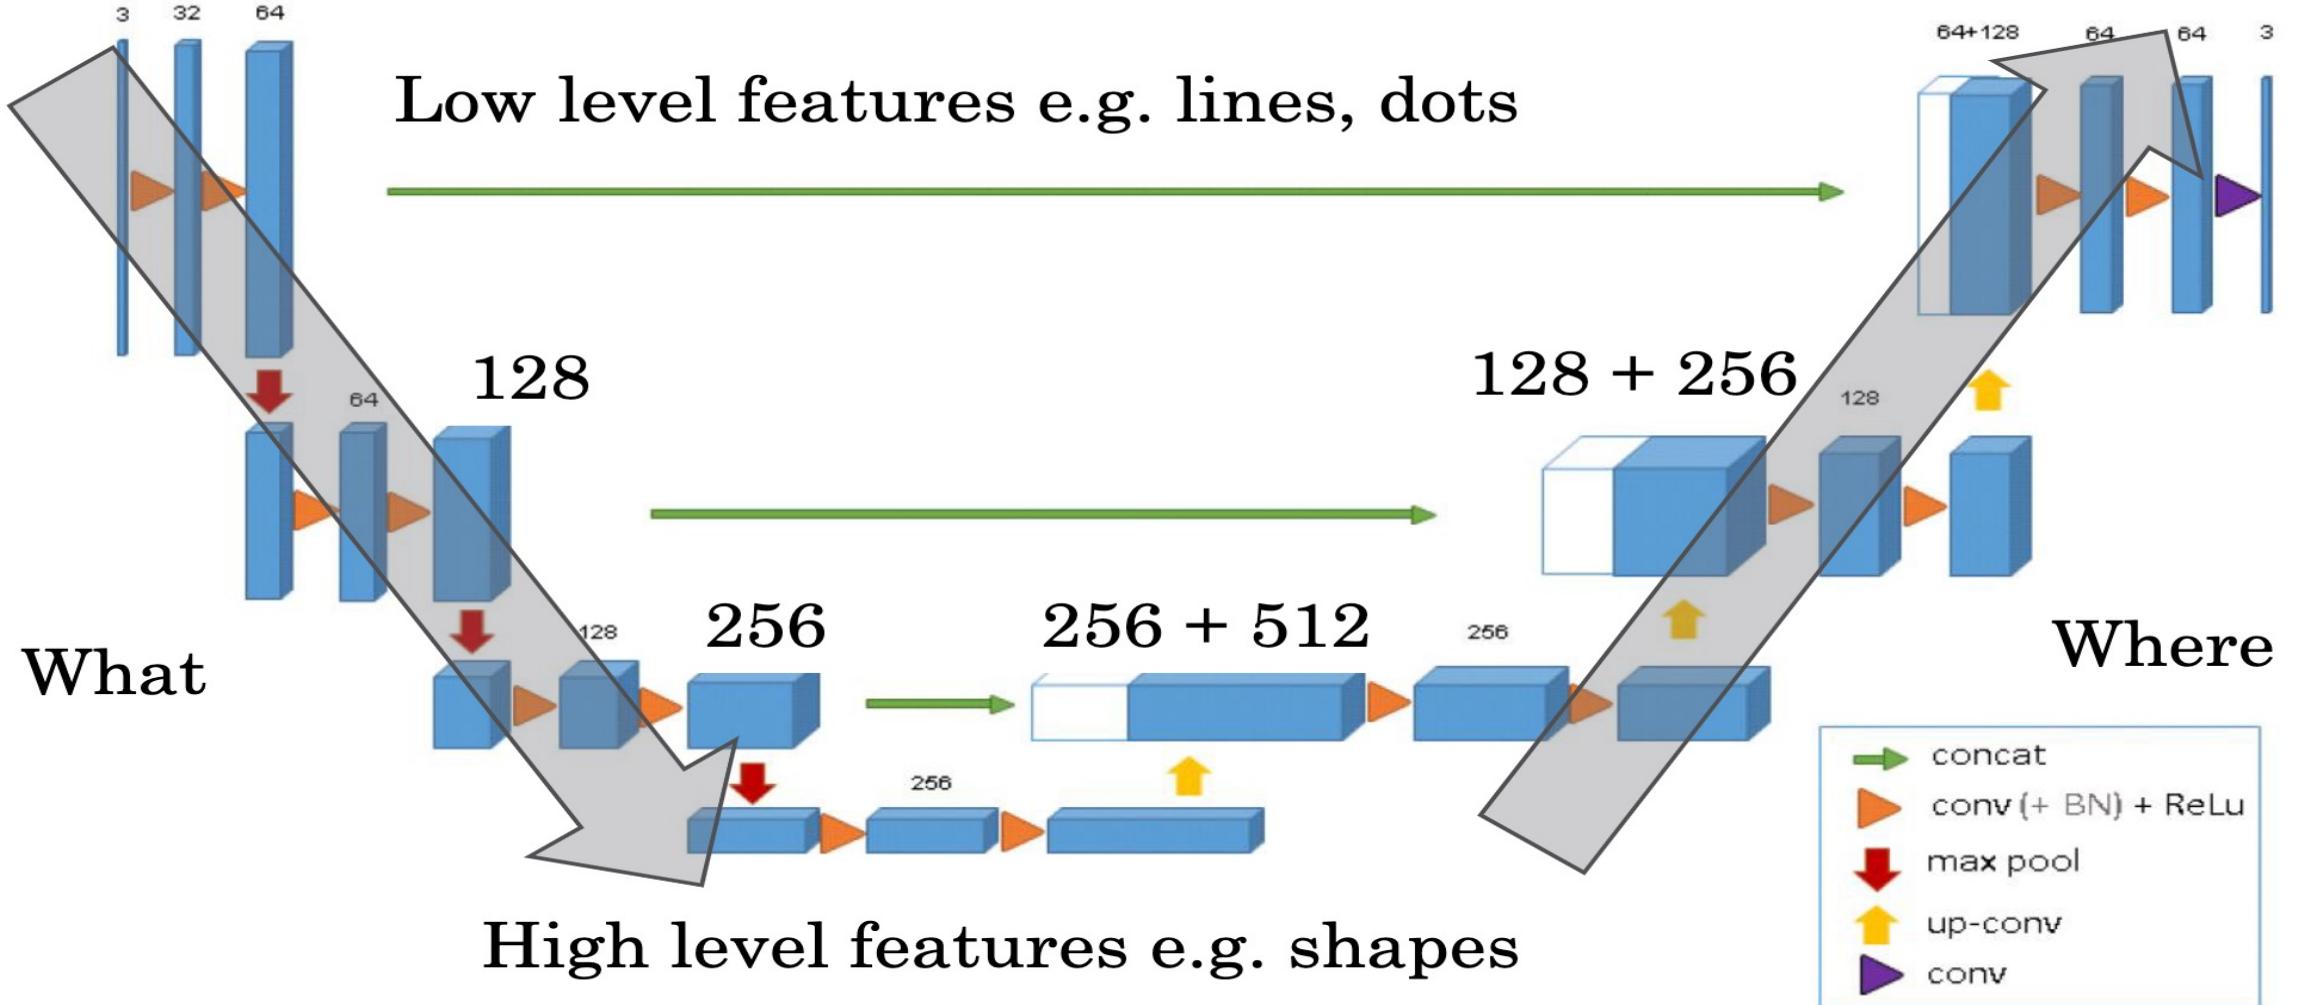
# Saved Models

- Weights:  
model.ckpt.data-00000-of-00001
- Weight Index:
  - model.ckpt.index
- Metadata of tensors
  - Model.ckpt.meta



# Evaluation Script

- `python evaluate.py --visu`

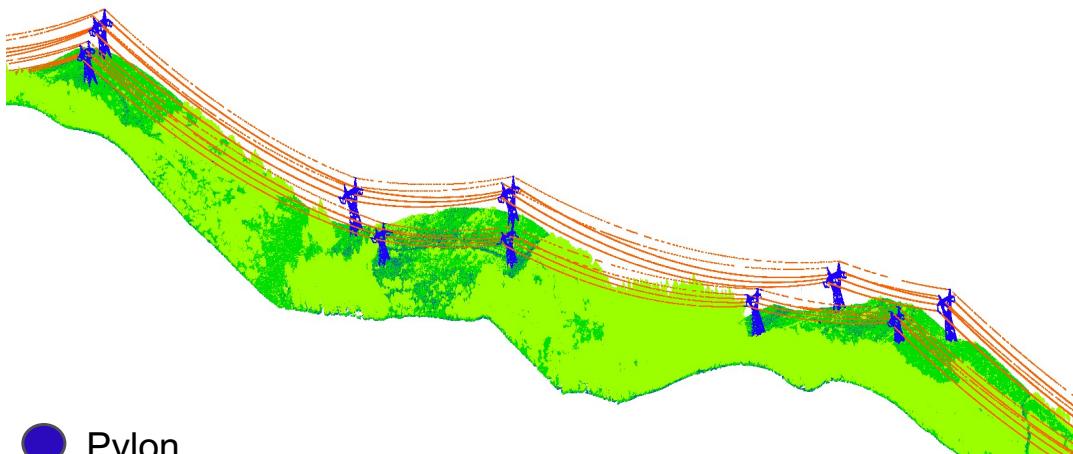


## 3D U-Net: Semantic Segmentation Demo

3D U-Net: Learning Dense Volumetric Segmentation from Sparse Annotation [Çiçek., et al. 2016]

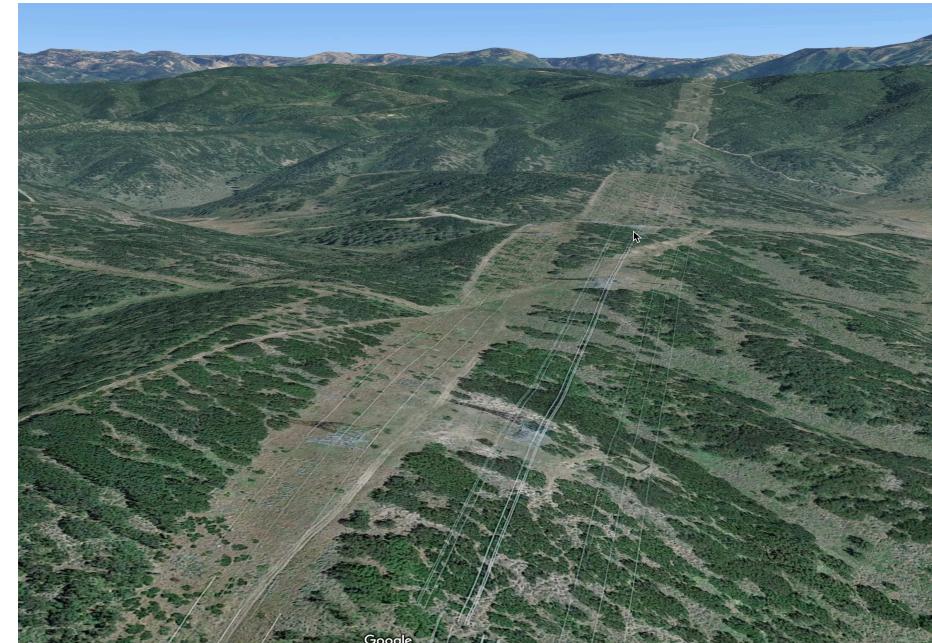
# Dataset

- Coverage area:  $66.7 \text{ km}^2$ ,
- Each scene area  $\cong 2.9 \text{ km}^2$ , total no. points  $\cong 2\text{m}$
- Sensor: Riegl Q560 laser scanner



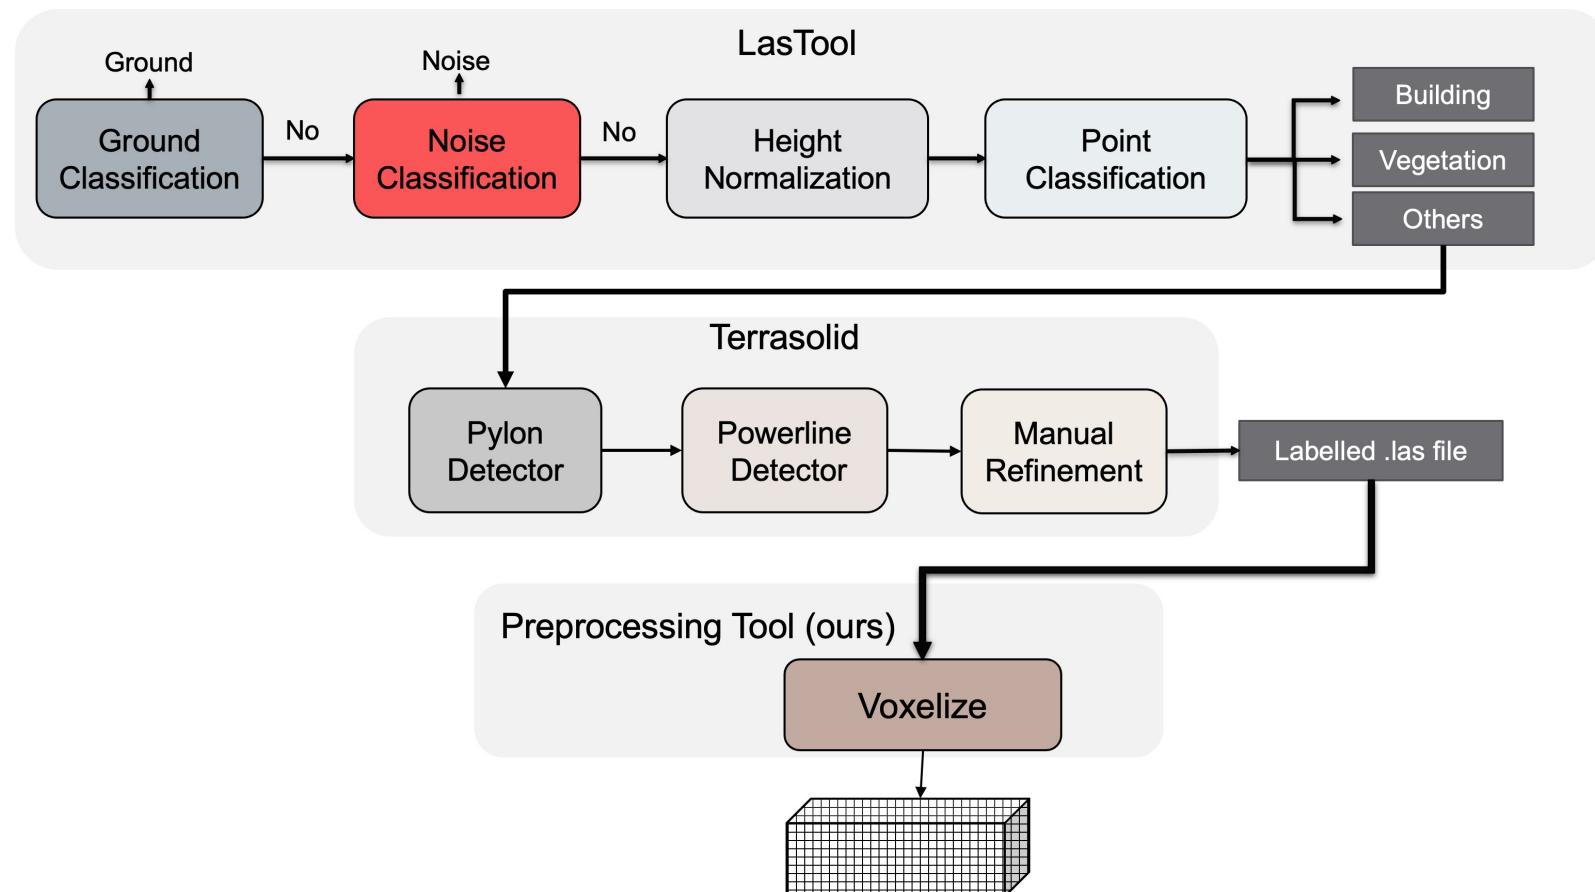
- Pylon
- Ground
- Trees
- Powerlines

Labelled Point cloud



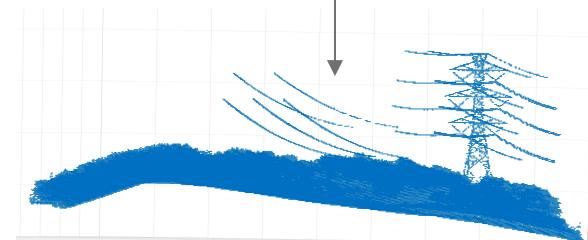
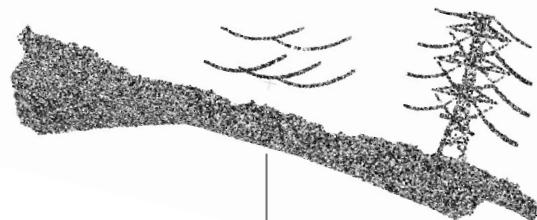
Steamboat Springs, Colorado, USA

# Preprocessing Pipeline



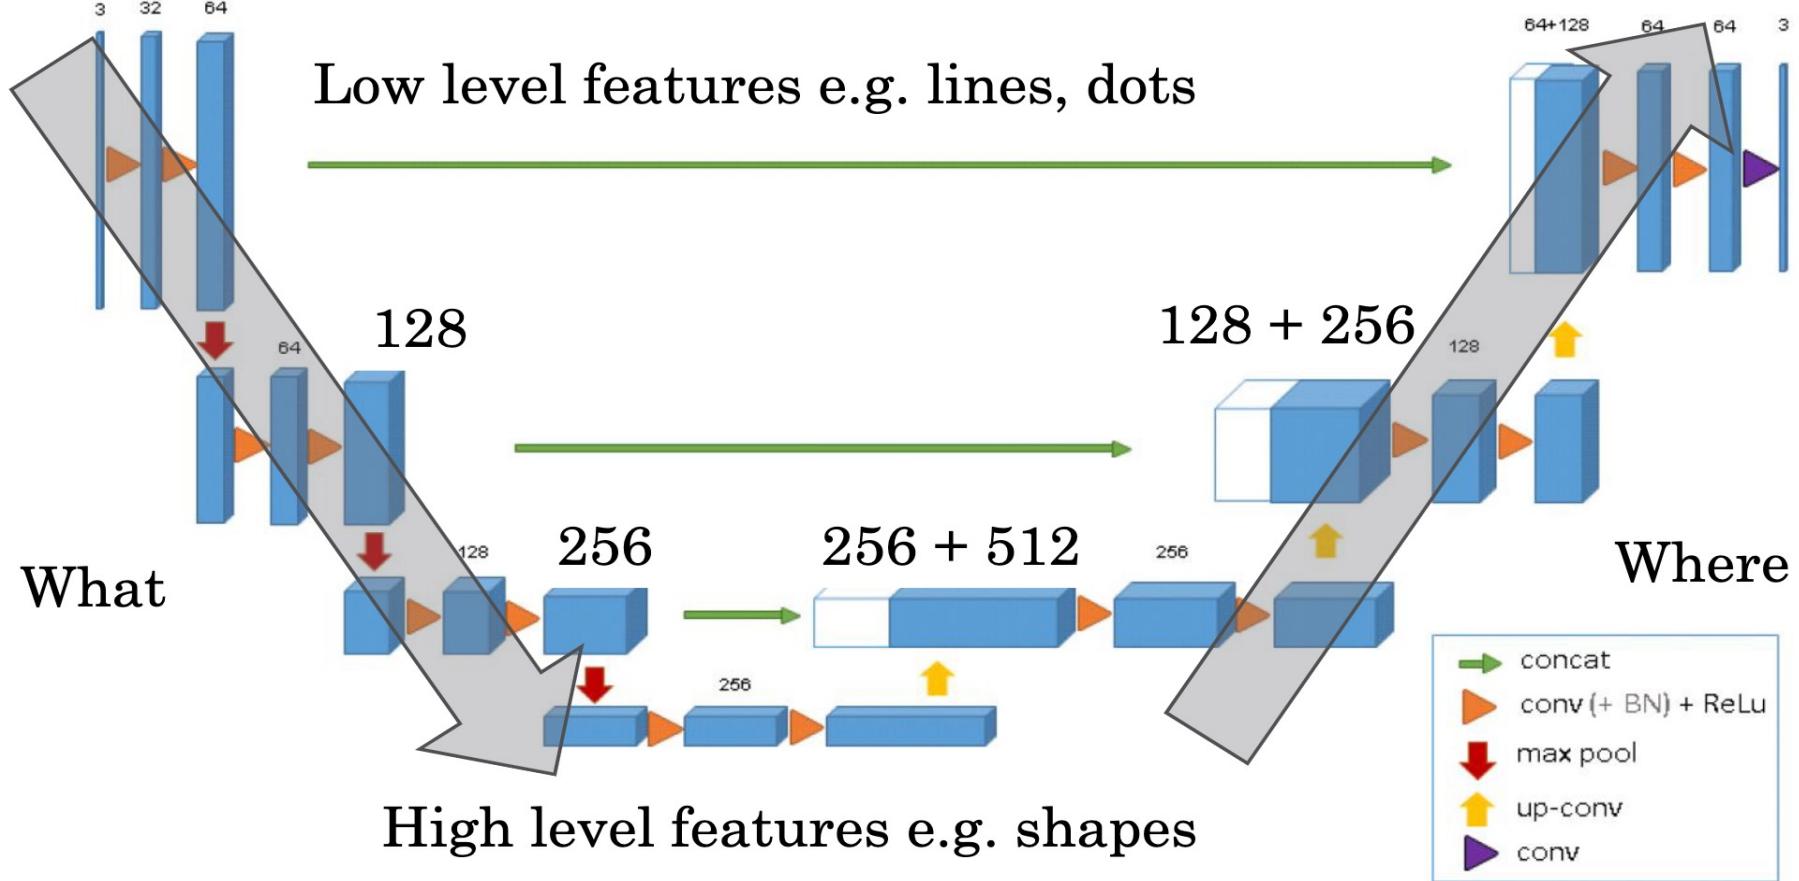
# Preprocessing Tool (ours)

**Input: Raw 3D point cloud**

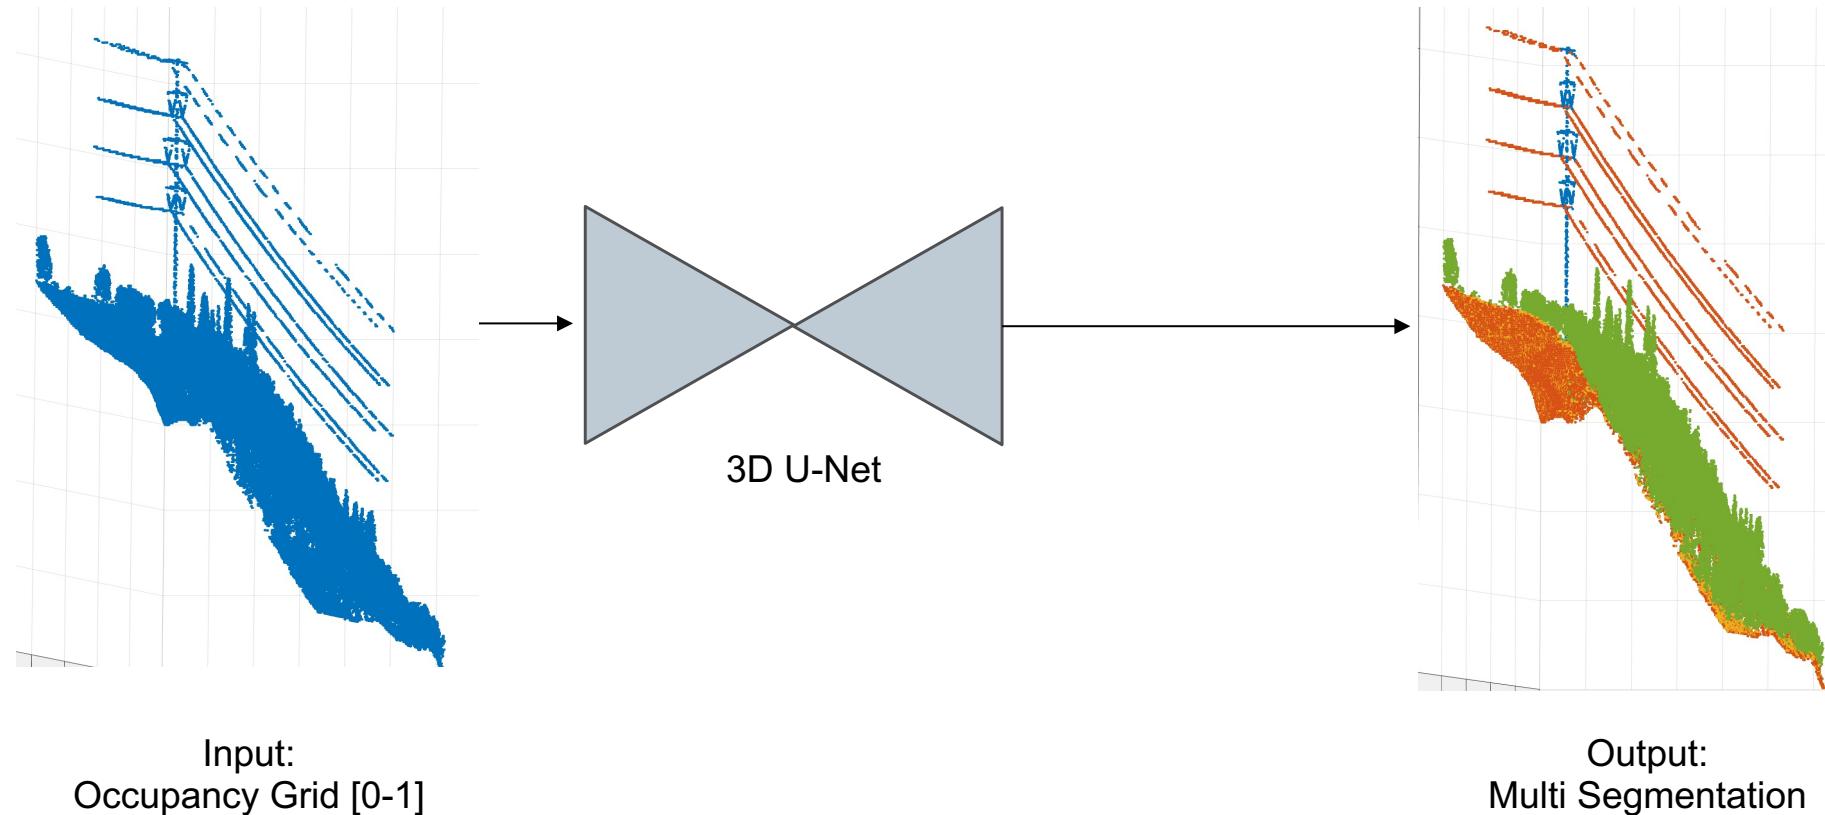


**Output: Normalized voxel grid**

# Network



# Semantic Segmentation



# Experimental Configuration

- Incremental K-Cross Validation Training
- Voxel size=  $1\text{m}^3$ ,
- Input grid size =  $32 \times 32 \times 32$  and Batch size=4 grids
- 2 GPU RTX 6000
- Training Time: 48-60 hours
- Inference: 2-3 mins

# Training

- python3 train.py
- Hyperparameter:

```
parser = argparse.ArgumentParser()
parser.add_argument('--gpu', type=int, default=0, help='GPU to use [default: GPU 0]')
parser.add_argument('--model', default='model', help='Model name [default: model]')
parser.add_argument('--category', default=None, help='Which single class to train on [default: None]')
parser.add_argument('--log_dir', default='log', help='Log dir [default: log]')
parser.add_argument('--XYZ', type=int, default=32, help='Point Number [default: 2048]')
parser.add_argument('--max_epoch', type=int, default=30, help='Epoch to run [default: 201]')
parser.add_argument('--batch_size', type=int, default=4, help='Batch Size during training [default: 32]')
parser.add_argument('--learning_rate', type=float, default=0.001, help='Initial learning rate [default: 0.001]')
parser.add_argument('--momentum', type=float, default=0.9, help='Initial learning rate [default: 0.9]')
parser.add_argument('--optimizer', default='adam', help='adam or momentum [default: adam]')
parser.add_argument('--decay_step', type=int, default=100000, help='Decay step for lr decay [default: 200000]')
parser.add_argument('--decay_rate', type=float, default=.8, help='Decay rate for lr decay [default: 0.7]')
parser.add_argument('--no_rotation', action='store_true', help='Disable random rotation during training.')

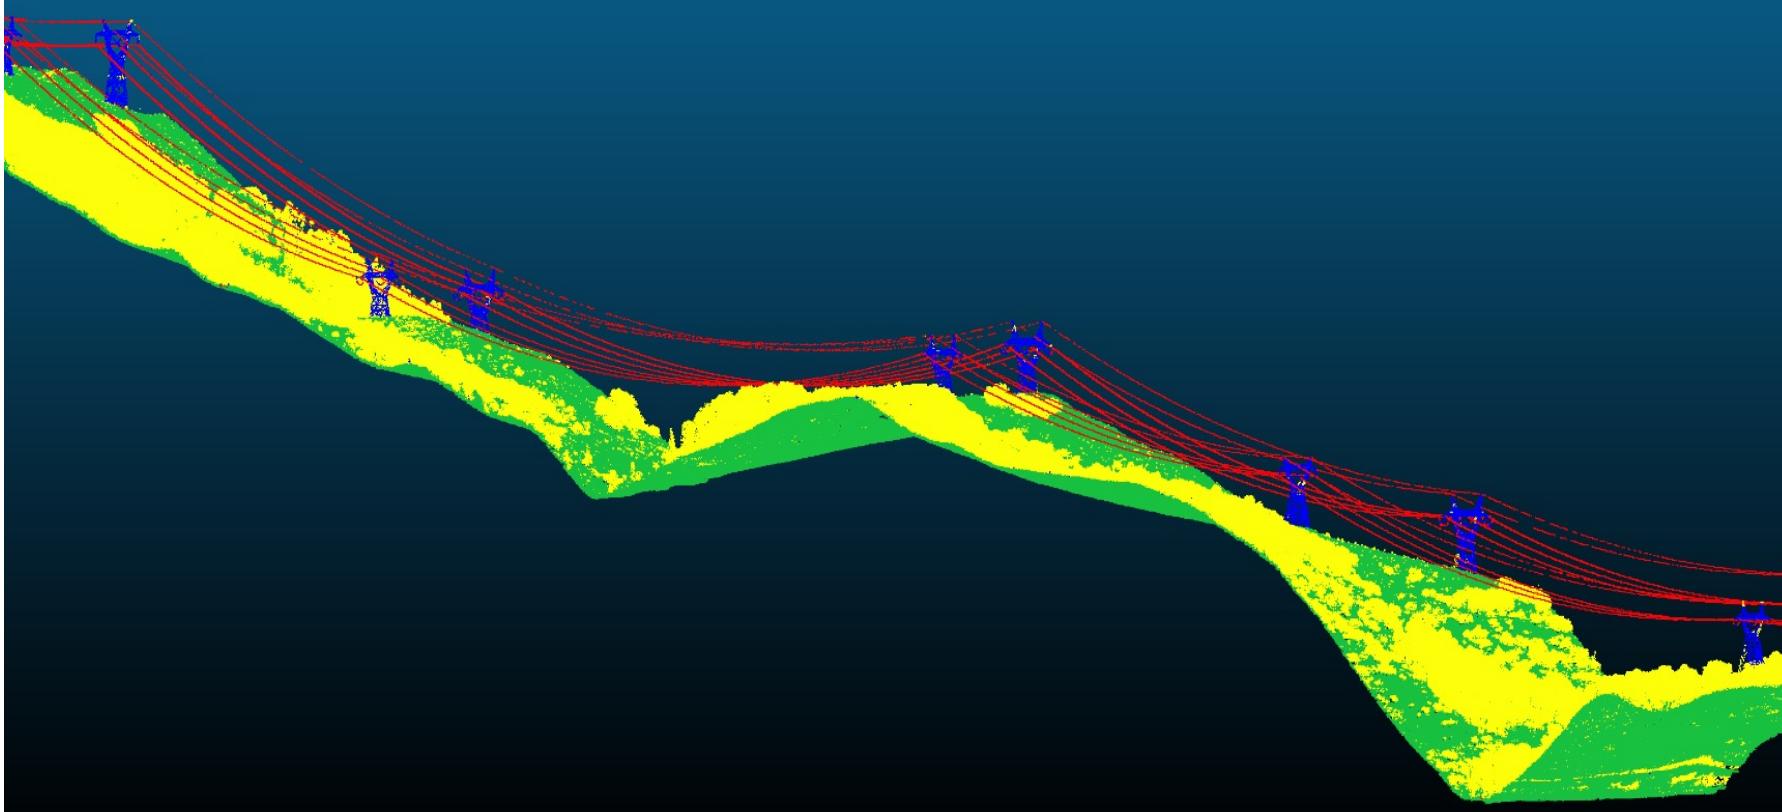
FLAGS = parser.parse_args()
```

# Encoder

# Decoder

```
main.py x Unet3D_network.py x
67     # Decoder
68     net = tf_util.conv3d_transpose(net, 16*base_fil, [2, 2, 2], padding=pad, stride=[2, 2, 2], bn=True,
69                                 is_training=is_training, scope='deconv1', bn_decay=bn_decay)
70     net_concat_3 = tf.concat([net_f_3, net], axis=-1)
71     net = tf_util.conv3d(net_concat_3, 8*base_fil, [3, 3, 3],
72                         padding=pad, stride=[1, 1, 1],
73                         bn=True, is_training=is_training,
74                         scope='conv8', bn_decay=bn_decay)
75     net = tf_util.conv3d(net, 8*base_fil, [3, 3, 3],
76                         padding=pad, stride=[1, 1, 1],
77                         bn=True, is_training=is_training,
78                         scope='conv9', bn_decay=bn_decay)
79     net = tf.layers.dropout(net, rate=drop_r, training=is_training)
80     net = tf_util.conv3d_transpose(net, 8*base_fil, [2, 2, 2],
81                                   padding=pad, stride=[2, 2, 2],
82                                   bn=True, is_training=is_training, scope='deconv2', bn_decay=bn_decay)
83     net_concat_2 = tf.concat([net_f_2, net], axis=-1)
84     net = tf_util.conv3d(net_concat_2, 4*base_fil, [3, 3, 3],
85                         padding=pad, stride=[1, 1, 1],
86                         bn=True, is_training=is_training,
87                         scope='conv10', bn_decay=bn_decay)
88     net = tf_util.conv3d(net, 4*base_fil, [3, 3, 3],
89                         padding=pad, stride=[1, 1, 1],
90                         bn=True, is_training=is_training,
91                         scope='conv11', bn_decay=bn_decay)
92     net = tf.layers.dropout(net, rate=drop_r, training=is_training)
93     net = tf_util.conv3d_transpose(net, 4*base_fil, [2, 2, 2],
94                                   padding=pad, stride=[2, 2, 2],
95                                   bn=True, is_training=is_training, scope='deconv3', bn_decay=bn_decay)
96     net_concat_1 = tf.concat([net_f_1, net], axis=-1)
97     net = tf_util.conv3d(net_concat_1, 2*base_fil, [3, 3, 3], padding=pad, stride=[1, 1, 1],
98                         bn=True, is_training=is_training,
99                         scope='conv12', bn_decay=bn_decay)
100    net = tf_util.conv3d(net, 2*base_fil, [3, 3, 3],
101                      padding=pad, stride=[1, 1, 1],
102                      bn=True, is_training=is_training,
103                      scope='conv13', bn_decay=bn_decay)
104    net = tf.layers.dropout(net, rate=drop_r, training=is_training)
105    net = tf_util.conv3d(net, 5, [1, 1, 1], padding=pad, stride=[1, 1, 1], bn=False, is_training=is_training,
106                        scope='conv14', bn_decay=bn_decay, activation_fn=None)
```

# Results



- Pylon
- Ground
- Trees
- Powerlines

# Conclusion

- There are pros and cons of point and voxel-based methods. Selection of architecture is based on problem domain and dataset.
- Tensorflow is rich api for deep learning programming.
- Hyperparameter tuning can help achieve optimal results
- Data preprocessing pipeline and experimental configurations are key steps.

# Acknowledgements

## 3D Mobile Mapping AI Project (Teledyne Optech and NSERC)



**NSERC**  
**CRSNG**



## Questions?